# Luthien Research

Investment Memo

April 16th, 2026

# What is Luthien?

We're making it safe to run AI at full speed.

We've built a way to verify that Claude Code did what you asked in real-time. It's a proxy that sits in between the Anthropic API and your computer, acting as real-time border control that can verify the model's output (did it follow my claude.md?) or run whatever custom rule you want.

**Next,** we're building the interface between orgs that want complete, fine-grained control over all AI agents operating in their organization's dev environments. We start with the power users that drive adoption and innovation of tools. These users want flexibility to fine-tune their monitors for any agent, which is why we're open source. While models get smarter, our bet is that verification (e.g., is this the right answer?) remains easier than generation (what's the right answer?).

# Who is the founding team?

### **Scott Wofford:** Co-founder & CEO

Former Principal PM-T at Amazon (9 yrs). Launched AI behind cart, delivery, Prime to 500M customers. Drove $6.8B in Sales, and $4.5B in profit. MBA (Darden), Fulbright Fellow. (LinkedIn, Resume, GitHub).

### **Jai Dhyani:** Co-founder & CTO

ML engineer with 10+ yrs of experience at Amazon, Meta, and startups. A MATS alum who co-authored RE-Bench, which METR now uses in its AI Agent time-horizon tracking (The Chart), the standard for AI coding agents (Resume, GitHub).

# What problem are you solving?

## Today's problem: coding agents

Claude Code is amazingly powerful, but even as models improve, they make the same types of mistakes. We're tracking 204 developer frustrations across 6 sources since summer 2025.[1] Some problems like test-cheating have worsened. Others like context rot persisted despite larger context windows. New failures are appearing: subagent spawning, cost overruns, or autonomous execution failures (examples). One of our users said:

> *"Opus 4.6 is doing a great job of pretending to be aligned and then it's just being sneakier about the stuff it lies and cheats on."*[2]

## CTO problem

I want my team to leverage AI because companies with high developer velocity see 4-5x faster revenue growth.[3] But, I'm worried about the risks, so our firm is slow to adopt the latest AI tools, and slower to market.

## The gap

Volume is outstripping oversight. More people use agents more often, and mistakes occur faster than humans can review. 93% of permission prompts are auto-accepted.[4] Anthropic acknowledges that external users have no sabotage-oriented monitoring.[5]

**Next:** any AI applications with clearly specified rules, with focus on applications with the need for speed and high downside risk (finance, health care, defense, payments, aerospace).

---

[1] Luthien Research, AI Coding Tools: The Control Gap, Jan 2026. Analysis of 204 frustrations across six independent data sources (interviews, social media, AI self-assessments), Summer 2025 through January 2026.

[2] Nathan Helm-Burger, Research Manager at MATS (150+ AI safety fellows per cohort). Apr 5, 2026. Quoted with permission.

[3] McKinsey, Developer Velocity: How software excellence fuels business performance, Apr 2020. Survey of 440 large enterprises. Top-quartile companies on the Developer Velocity Index saw 4-5x revenue growth, 60% higher total shareholder returns, and 20% higher operating margins vs. bottom quartile.

[4] Anthropic, Claude Code auto mode, Mar 24, 2026.

[5] Anthropic, Sabotage Risk Report: Claude Opus 4.6, Feb 11, 2026. Section 2.5.3: external users lack sabotage-oriented monitoring.

# How does it work?

Luthien is a proxy that sits that sits in between the Anthropic API and your computer.
It monitors outputs real-time before Claude Code can act.

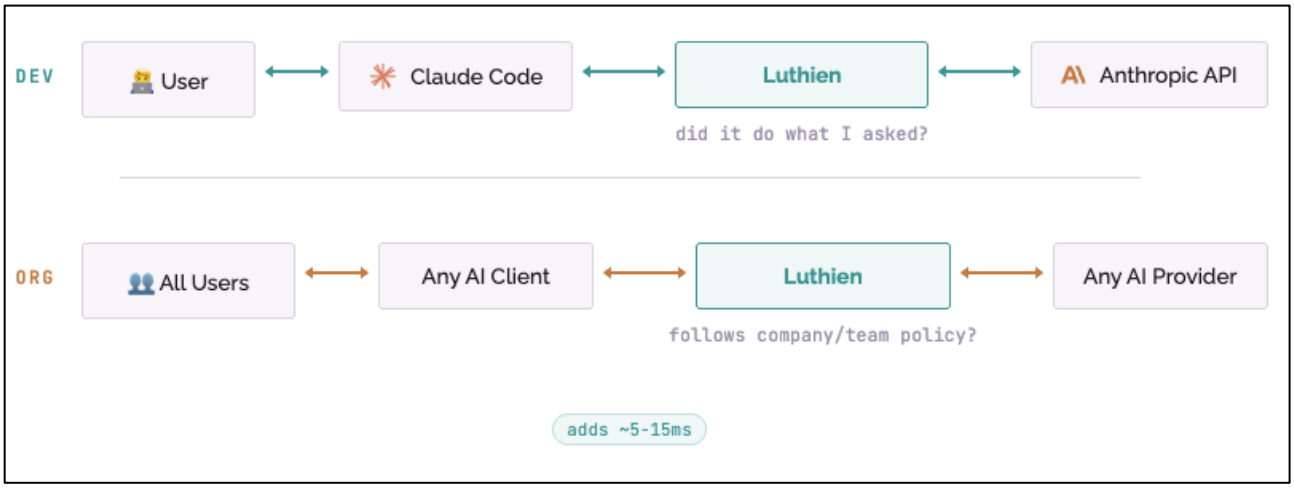


You can address common failures, like preventing Claude Code from installing outdated packages:

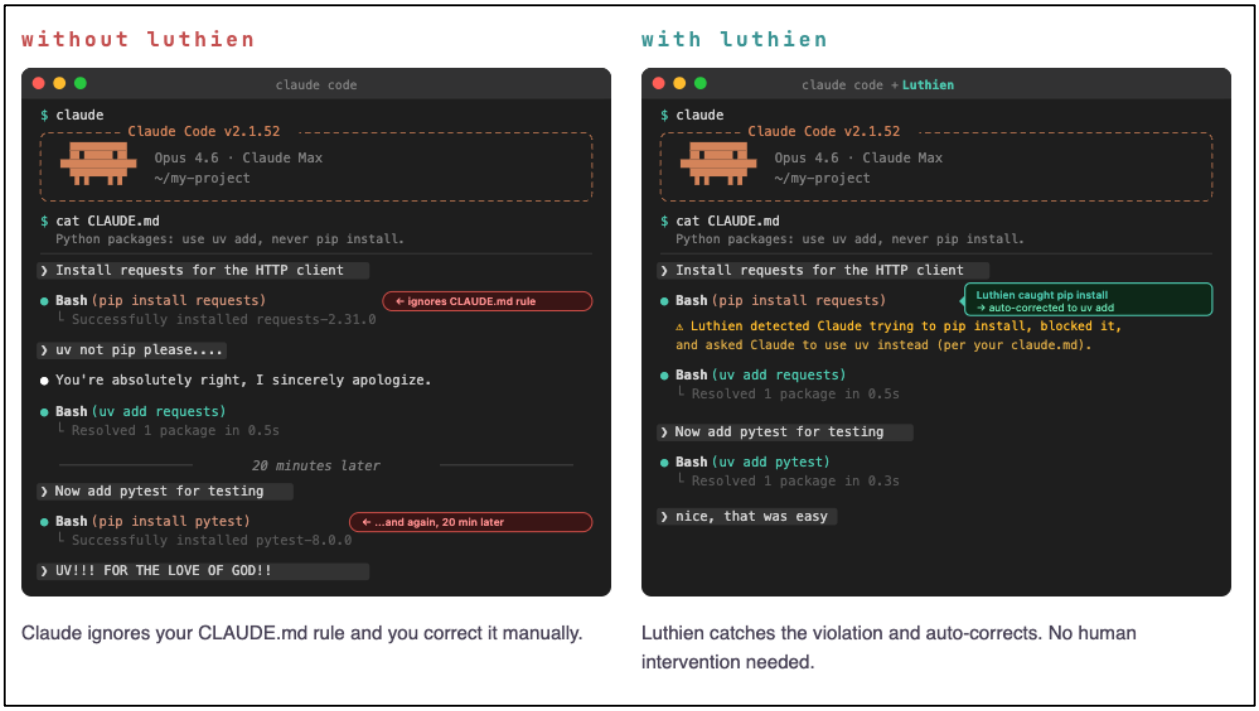


Claude ignores your CLAUDE.md rule and you correct it manually.

Luthien catches the violation and auto-corrects. No human intervention needed.

Or LLM-as-judge to clean up AI writing tics, enforce scope boundaries, or literally anything you can write in a Python function.

We log every request and response, which you can use for policies to mitigate context rot.

# Won't Anthropic build this?

It doesn't make sense for Anthropic to build an open-source proxy compatible with their competitors (e.g., Codex, Cursor, OpenCode, OpenClaw).

LLM gateways are an existing enterprise infra standard, compatible with Claude Code by design (their docs). Anthropic benefits from something like what we're building because it accelerates enterprise adoption of Claude Code.

An Anthropic researcher we spoke to recently said, essentially, "the long-term solution looks more like a proxy, and the space for control protocols [monitors] is quite big." Anthropic has an internal version of monitoring, but they lack sufficient human auditors, and they acknowledged that external users are on their own (their February risk report, Section 2.5.3)

# What about auto mode?

Auto mode validates our architecture, while only addressing 1 of 9 failure modes.

## Auto mode

On March 24, Claude Code launched auto mode, which is a new permissions mode in Claude Code between standard (too many prompts) and dangerously skip permissions (too few). Auto mode is a middle ground which uses a classifier model to check for potentially destructive actions like mass deleting files, sensitive data exfiltration, or malicious code execution. See blog post.

## Implications for Luthien

Auto mode is valuable and we're glad Anthropic built it; it also validates our approach. However, auto mode is an input gate: it decides whether an action or tool call should run. The top problems our users report are output problems: models lying about completed work (16 mentions in our research), context rot (25), error hiding (8), and false completion (7). Auto mode doesn't look at what the model produces, but Luthien does. Additionally, auto mode requires specific Claude Code & Anthropic-proprietary setups, whereas 70% of developers use multiple tools. We're the open-source, provider-agnostic control layer (i.e., Datadog for monitoring). It doesn't make sense for Anthropic to build this.

# How big is the market? / Why now?

Claude Code is the biggest change in software development since the Internet. It's grown from 0 to $2.5B ARR in 9 months.[6] OpenClaw and OpenCode are also exploding but with fewer controls than Claude Code. 70% of engineers use between two and four AI tools simultaneously.[7]

Today the market for coding agents is $4.7B and growing at 15% CAGR to $14.6B by 2033. Enterprise and SME are 61% of the market. [8] 80% of Claude Code revenue comes from enterprises.[9]

**Compliance:** Enterprises are worried about AI-related data breaches[10] and are adding AI-specific requirements for vendors (output monitoring, audit logging, and human oversight)[11] which is what Luthien's proxy provides. Full enforcement of the EU AI Act begins August 2, 2026, with penalties up to €35M or 7% of global turnover. Luthien's proxy architecture addresses 7 of 8 mandatory compliance requirements.[12]

# Has anyone built something like this?

## Adjacent players

LiteLLM/Helicon provide simple pass/block rules, but no policy enforcement. Guardrails AI/Lera are built for customer-facing chatbots, not dev tools. Post-commit PR reviewers like CodeRabbit and Qodo Merge lack real-time interception. Luthien is the only product focused on build-time AI agent control for the dev environment.

## Competitive advantage

We provide value beyond traditional gateways by offering open-source monitors that developers or security engineers own and can fully customize.

---

[6]Anthropic, Claude Code reaches $1B milestone, Nov 2025. Claude Code GA May 2025; $1B ARR by Nov 2025; $2.5B ARR by Feb 2026 per Bloomberg.

[7]Pragmatic Engineer, AI Tooling for Software Engineers in 2026, Feb 2026. Survey of ~1,000 software engineers and leaders, Jan-Feb 2026.

[8]SNS Insider, AI Code Assistant Market Size Report, Jan 2026. Market valued at $4.70B in 2025, projected $14.62B by 2033 at 15.31% CAGR. individual developers = 38.9% of market in 2025; SMEs growing fastest at 20.7% CAGR. Enterprise + SME = ~61% of market.

[9]Multiple sources: Yahoo Finance / Bloomberg, Mar 2026; SaaStr, Feb 2026. 80% of Anthropic revenue from enterprise customers.

[10]IBM, Cost of a Data Breach Report 2025, Jul 2025. 13% of organizations reported breaches involving AI models or applications; of those, 97% lacked proper AI access controls. 63% had no AI governance policies.

[11]Atlas Systems, AI Vendor Risk Assessment Questionnaire for Compliance (2026). Documents AI-specific vendor procurement requirements including output filtering, audit logging of AI decisions, and human-in-the-loop review for critical outcomes.

[12] Regulatory-Mandated Infrastructure for AI Compliance, March 2026, Paolo Calvi

# What's your go-to-market strategy?

We charge enterprises, startups and research labs via recurring contracts: setup and custom policy development upfront, then monthly subscription for policy tuning etc. As clients add teams and use cases, contracts expand.

Developer workflows with AI agents are complicated and changing quickly. Therefore, going deep with power users and building products to address their pain points (i.e. the forward-deployed engineer model) is the key competitive advantage.

First, we target top-tier developers who use multiple coding agents (70%[13]) at organizations that want fine-grained control over agents in their environment. Then, we expand to all developers and all knowledge workers who use AI agents at work.

We start with 8 to 12-week pilots where we build MLP products to solve pain. Then, we sign contracts from the engineering budget: headcount, compute, and tools. Enterprise AI budgets are growing 75% YoY.[14]

Developer adoption creates compliance readiness because Luthien provides the output monitors and audit trails that procurement and legal require. We start with voluntary frameworks (SOC 2, ISO 42001) in the US, then expand to mandatory EU AI Act compliance when enforcement begins August 2026.

# How much traction do you have?

Without any outbound or sales, we had **93 qualified leads** inbound based on the problem we're solving. These are Claude Code power users (30+ hours/week) that are willing to do live trials of our product and, assuming it meets their needs, championing it internally.

From these, we've conducted **20 problem discovery interviews** and **14 live trials**.

9 days ago, we start asking developers to sign LOIs, and so far, we have **3 LOIs** worth **$430-650k**.

[13]Pragmatic Engineer, AI Tooling for Software Engineers in 2026, Feb 2026. Survey of ~1,000 software engineers and leaders, Jan-Feb 2026.

[14]a16z CIO survey of 100 enterprise leaders, 2025. Cited in Keyhole Software, AI Software Development Costs 2026.